




CIFAR-10 Classification with CNNs

JOSHUA LABRANCHE

JUNGHWAN KIM

JONATHAN LAPHAM

Outline



- ▶ **Introduction**
- ▶ Dataset
- ▶ Architecture
- ▶ Results
- ▶ Conclusions


Introduction

- **Constructed original CNN using inspiration from AlexNet, VGG, and ResNet architectures**
- **Architecture design focused on improving a baseline model with regularization and normalization methods.**
- **Computational cost was a factor in architecture design, as portability was a desired characteristic**
- **Lowest error rate achieved: 13.2%**

What Did We Do?

Outline

▶ Introduction



▶ Dataset

▶ Architecture

▶ Results

▶ Conclusions


Dataset

- Architecture was trained and tested using the CIFAR-10 dataset
 - Dataset consists of 60,000 32x32x3 images categorized into 10 categories:
 - airplanes
 - cars
 - birds
 - cats
 - deer
 - dogs
 - frogs
 - horses
 - ships
 - trucks.
- Part of the larger “80 Million Tiny Images” data set found at:
<http://groups.csail.mit.edu/vision/TinyImages/>

A. Krizhevsky, “Learning multiple layers of features from tiny images,”
tech. rep., 2009.

Outline

- ▶ **Introduction**
- ▶ Dataset
- ▶ Architecture
- ▶ Results
- ▶ Conclusions



Architectures

Untuned Model A	Untuned Model B	Untuned Model C	Untuned Model D
2D Conv (256, (3x3)) 2D Conv (256, (3x3)) 2D Conv (256, (3x3)) MaxPool((2x2), stride=2)	2D Conv (64, (3x3)) 2D Conv (64, (3x3)) 2D Conv (64, (3x3)) MaxPool((2x2), stride=2)	2D Conv (256, (3x3)) 2D Conv (256, (3x3)) 2D Conv (256, (3x3)) MaxPool((2x2), stride=2)	2D Conv (256, (3x3)) 2D Conv (256, (3x3)) MaxPool((2x2), stride=2)
2D Conv (128, (3x3)) 2D Conv (128, (3x3)) MaxPool((2x2), stride=2)	2D Conv (64, (3x3)) 2D Conv (64, (3x3)) MaxPool((2x2), stride=1)	2D Conv (256, (3x3)) 2D Conv (256, (3x3)) MaxPool((2x2), stride=2)	2D Conv (256, (3x3)) MaxPool((2x2), stride=1)
2D Conv (64, (3x3)) 2D Conv (64, (3x3)) MaxPool((2x2), stride=2)	2D Conv (32, (3x3)) MaxPool((2x2), stride=2)	2D Conv (256, (1x1)) MaxPool((2x2), stride=1)	2D Conv (256, (1x1)) MaxPool((2x2), stride=1)
NO LAYER	2D Conv (64, (3x3)) 2D Conv (64, (3x3)) MaxPool((2x2), stride=1)	2D Conv (64, (3x3)) MaxPool((2x2), stride=1)	2D Conv (64, (3x3)) MaxPool((2x2), stride=1)
NO LAYER	NO LAYER	2D Conv (256, (1x1)) MaxPool((2x2), stride=2)	2D Conv (256, (1x1)) MaxPool((2x2), stride=2)
Flatten()	Flatten()	Flatten()	Flatten()
Dense(512)	Dense(512)	Dense(512)	Dense(512)
Dense(512)	Dense(512)	Dense(512)	Dense(512)
Dense(10)	Dense(10)	Dense(10)	Dense(10)

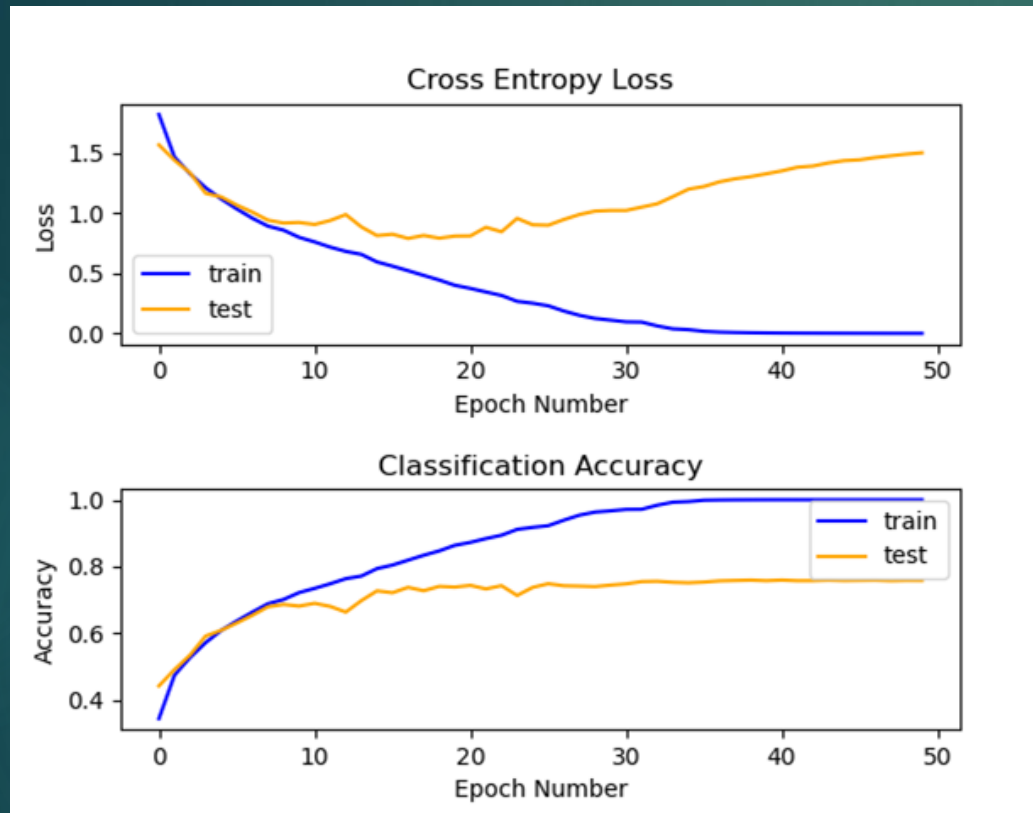
Outline

- ▶ Introduction
- ▶ Dataset
- ▶ Architecture
- ▶ Results
- ▶ Conclusions

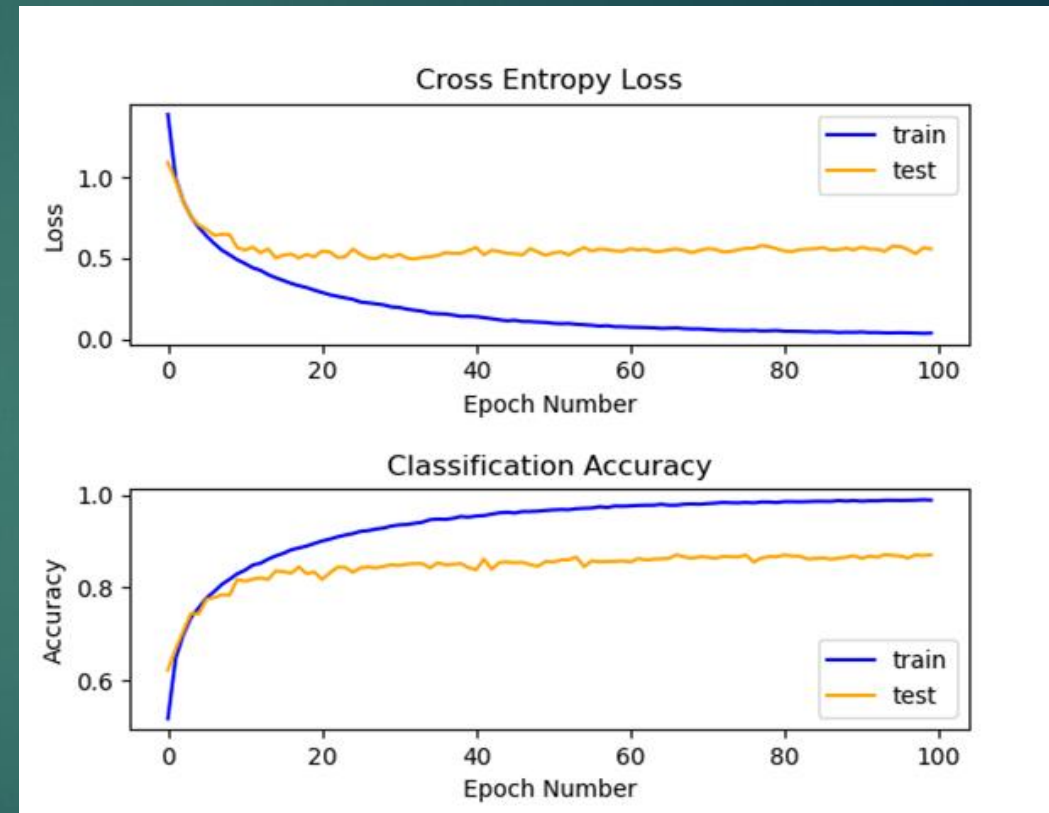


Results

Baseline Model D



Model D with Data Augmentation, Batch Normalization, and a Reduced Batch Size (64)



Error Rate Minimum: 13.2%

Outline

- ▶ Introduction
- ▶ Dataset
- ▶ Architecture
- ▶ Results
- ▶ Conclusions



Conclusion

- Model D was the best performing of our baseline models.
- We tried to improve performance by adding techniques such as weight decay, drop-out, scaling drop-out, data augmentation, bottlenecking, and batch normalization.
- The best measured performance was the baseline model with the addition of data augmentation and batch normalization, along with a reduced batch size
- Performing only slightly worse was Model D with the addition of data augmentation and constant drop-out with an error rate of 14.8%

Data Augmentation and Batch Normalization!